

BYZANTINE GENERALS PROBLEM

GRAHAM SWAIN

ABSTRACT. In this paper we explore the Byzantine Generals problem; an abstraction used to discuss the reliability of computing systems.

1. INTRODUCTION

The Byzantine Generals Problem is a classic example that is used to abstract discussions of computing systems. For a computing system to be reliable it must be able to deal with one or more of its components failing. In the Byzantine Generals Problem each component is represented as a general, the components that failed are traitorous generals. We will mostly be exploring when all the components, or generals, have a direct line of communication, but we will briefly discuss about when that is not the case.

The basis of the problem is we imagine that there are multiple divisions of the Byzantine army sieging a city, each division led by a general. The generals are only able to communicate by messengers. The general needs to decide and agree on a plan. Unfortunately, some of the generals could be traitors, trying to prevent the loyal generals from reaching an agreement. So the generals need to have an algorithm to guarantee that:

- A. All loyal generals decide upon the same plan of action.
- B. A small number of traitors cannot cause the loyal generals to adopt a bad plan.

The goal is to ensure that the loyal generals are able to agree on a reasonable plan, regardless of traitorous generals. It is hard to quantify exactly what a “bad” plan is, but we do not have to for our algorithm. How the generals reach a decision is more important. Let there be n generals and v_i is the information communicated by the i^{th} general. Each general has to use some method of taking all of the plans communicated to them, v_1, \dots, v_n , and make a decision based off of them. Having all of the generals use the same method of decision making would achieve Condition A.

Condition B can be met by using a robust decision making method. To simplify the problem we say that the message can only be one of two orders “ATTACK” or “RETREAT”, making v_i General’s i opinion on which is the best option. Each general can take a majority vote of v_1, \dots, v_n and let whichever option wins be the “best” decision. This method only lets a small number of traitors affect the result if the loyal generals are split fairly evenly between the two options, in which case neither option could be deemed “bad”.

To help ensure that Condition A and Condition B are met, we need to meet two other conditions:

1. Every loyal general must obtain the same information v_1, \dots, v_n .
2. If the i^{th} general is loyal, then the value that they sends must be used by every loyal general as the value of v_i .

Condition 1 could also be written as:

- 1'. For every i , any two loyal generals use the same value of v_i .

These are important because a loyal general cannot take a value v_i at face value, because a traitorous general could send different values of v_i to different generals. Even though the generals cannot necessarily trust every message, we must find a way to guarantee that loyal generals use the values sent to them by other loyal generals.

Since Condition 1' and Condition 2 are both contingent on a single v_i sent by the i^{th} general, we can focus on how a single general sends values to others. We phrase this problem by referring to a commander sending messages to their lieutenants.

Byzantine Generals Problem. A Commanding general must send an order to their $n - 1$ lieutenants such that:

- IC1. All loyal lieutenants obey the same order.
- IC2. If the commander is loyal, then every loyal lieutenant obeys the order they sends.

Conditions IC1 and Condition IC2 are called the interactive consistency conditions. If the commander is loyal, then IC1 follows IC2. However, the commander can be a traitor.

2. THREE GENERALS

When sending only oral messages, at least two thirds of the lieutenants must be loyal for any solution to work. The distinction of oral message is important, as it means that the contents of a message is completely up to sender. So a traitorous sender can send any message they want.

We will look at the simplest form of this problem, where there are three generals (one commander and two lieutenants). There are two possible orders: "ATTACK" or "RETREAT". No solution can exist for this if there is even a single traitor.

Situation 1. The commander is loyal, but one of the lieutenants is a traitor.

The commander sends out the command ATTACK to both lieutenants. Both lieutenants rely the order they receive to the other, but since one lieutenant is a traitor, they rely the message RETREAT. The loyal lieutenant receives the set of orders: ATTACK, RETREAT, there is no way for him to decide which is the correct order.




FIGURE 1. Three General Problem with a loyal Commander.

Situation 2. The commander is not loyal, both lieutenants are.

The commander sends **ATTACK** to Lieutenant 1 and **RETREAT** to Lieutenant 2. Since they are both loyal they rely the commands to each other. Lieutenant 1 receives the set of orders: **ATTACK**, **RETREAT**.

Since Lieutenant 1 receives the set of orders: **ATTACK**, **RETREAT** in both situations, there is no way for them to know who they traitor is, so they follows their commander's orders to attack in both situations, regardless of if the commander is a traitor.

We can make a similar argument that Lieutenant 2 must follow the **RETREAT** command in the second situation. This violates IC1 since the two loyal lieutenants obey different orders.

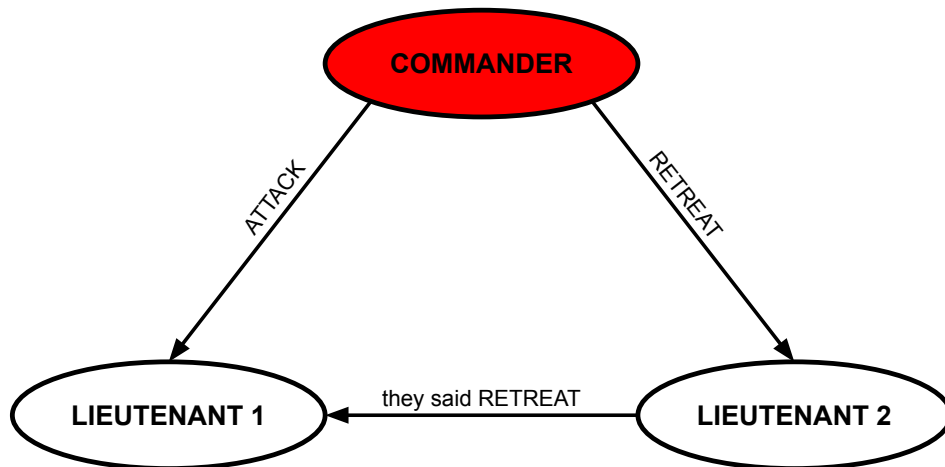


FIGURE 2. Three General Problem with a disloyal Commander.

The authors have a formal solution for the three generals problem in their 1980 paper "Reaching Agreement in the Presence of Faults" if you want to read more about it. For the purposes of this paper it is only important to know that no solution exists for three generals if there is a single traitor.

Using the result that a single traitor can disrupt the the three generals problem, we can show that no solution exists with less than $3m + 1$ generals and m traitors. To help differentiate between the problems, the generals from the three generals problem will be referred to as Byzantine generals and the generals from the assumed solution will be referred to as Albanian generals.

Proof. We want to prove that no solution exists with $n < 3m + 1$ generals with m traitors and $n > 3$. We will prove this by contradiction.

Assume that a solution exists for $3m$ or less Albanian generals. We will show a solution exists for three Byzantine generals with a single traitor. Each Byzantine general represents at most m Albanian generals. Since the Albanian commander needs to be represented as well, the Byzantine commander represents the Albanian commander as well as at most $m - 1$ Albanian lieutenants. We know that there is a single Byzantine traitor. Since each Byzantine general represents at most m Albanian generals, we know there is at most m Albanian traitors. The assumed solution means that IC1 and IC2 is true for the Albanian generals. Since each Albanian general is represented by a Byzantine general, then IC1 and IC2 must also be true for the Byzantine generals, which we know is impossible, forming a contradiction. \square

3. ORAL SOLUTION

The above section shows that we need at least $3m + 1$ generals to solve the Byzantine Generals Problem using oral messages, with m traitors. In this section we will create an algorithm that will work for $3m + 1$ or more generals. We make three assumptions to define an oral message:

- A1. Every message that is sent is delivered correctly.
- A2. The receiver of a message knows who sent it.
- A3. The absence of a message can be detected.

These assumptions each help prevent foul play by traitorous generals. A1 prevents traitor from stopping messages from being received. A2 make it to where a traitor cannot impersonate another. A3 stops traitors who try and prevent a decision from being made by not sending messages.

For this problem we assume that every general is able to send a message to every other general. That is often not the case and will be discussed in a later section. We also assume that every loyal general completes their algorithm. In the case of a traitorous general who does not send any orders at all, we assume that lieutenants default to the RETREAT order.

We will define an algorithm for oral messages $OM(m)$, for all nonnegative integers m . In this algorithm a commander sends an order to $n - 1$ lieutenants. $OM(m)$ will solve the Byzantine Generals Problem for $3m + 1$ more generals when there is at most m traitors.

As stated above, when a general receives a message, or a value, from a general, we denote it as v_i for the i^{th} general. Whenever a general has a set of all the messages, v_i, \dots, v_n , they use a majority function $\text{majority}(v_i, \dots, v_n) = v$, v being the decision the lieutenant comes too. Two choices for deciding the majority listed in the paper are:

1. The majority value among the v_i (the value that occurs the most) if one exists, or default to the value RETREAT.
2. The median value of the v_i , if the set is ordered.

The algorithm is recursive, so starts with a base case of OM(0) and uses the majority function.

Algorithm OM(0)

1. The commander sends their value to every lieutenant.
2. Each lieutenant uses the value they received from the commander. If they received no value, default to RETREAT.

Algorithm OM(m), $m > 0$

1. The commander sends his value to every lieutenant.
2. For each i ,
 - a. Lieutenant i receives a value v_i from the commander. Default to RETREAT if they receive no value.
 - b. Lieutenant i acts as the commander in OM($m - 1$) to send the message to each of the remaining $n - 2$ lieutenants.
3. For each i , and each j not equal to i ,
 - a. let v_j be the value Lieutenant i received from Lieutenant j in step (2b). Default to RETREAT if Lieutenant i received no value from Lieutenant j .
 - b. Lieutenant i uses the value $\text{majority}(v_1, \dots, v_{n-1})$.

Let's look at an example with four generals and one traitor, or $n = 4$ and $m = 1$. There are two situations, where one of the lieutenants is a traitor or where the commander is a traitor.

Situation 1. Lieutenant 3 is a traitor.

In the first step of the algorithm OM(1), the commander sends the order x to all three lieutenants. In the second step, Lieutenant 1 sends the value x to Lieutenant 2 using OM(0). Lieutenant 3 sends y to Lieutenant 2. In step 3, Lieutenant 2 has the set of orders $v_1 = v_2 = x$ and $v_3 = y$. Using $\text{majority}(x, x, y) = x$, Lieutenant 2 follows the order x sent to him by the commander. We can use a similar process to find that Lieutenant 1 also follows order x . So both loyal lieutenants followed the order of their loyal commander. This satisfies IC1 and IC2.

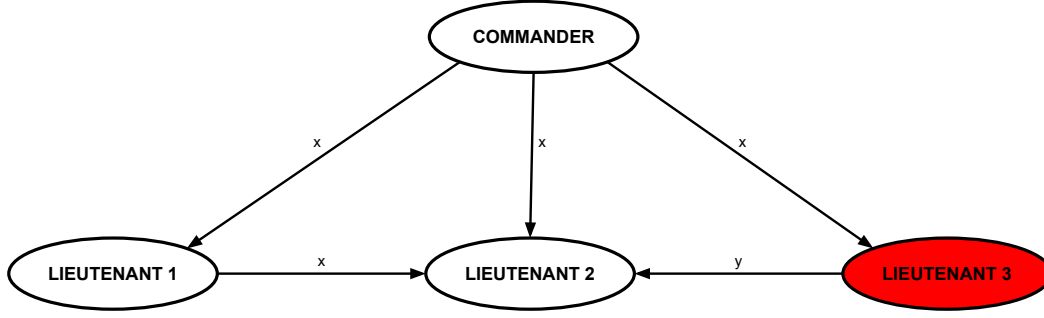


FIGURE 3. OM(1) with a disloyal lieutenant.

Situation 2. The commander is a traitor.

The commander sends the arbitrary orders x, y, z to Lieutenants 1, 2, 3 respectively. In step 2 each lieutenant sends the order they got to the others. In step 3 each lieutenant ends up with $v_1 = x, v_2 = y, v_3 = z$ and the majority function $majority(x, y, z)$. So they each will follow the same order regardless of if the values of x, y, z are equal to one another or not. This satisfies IC1.

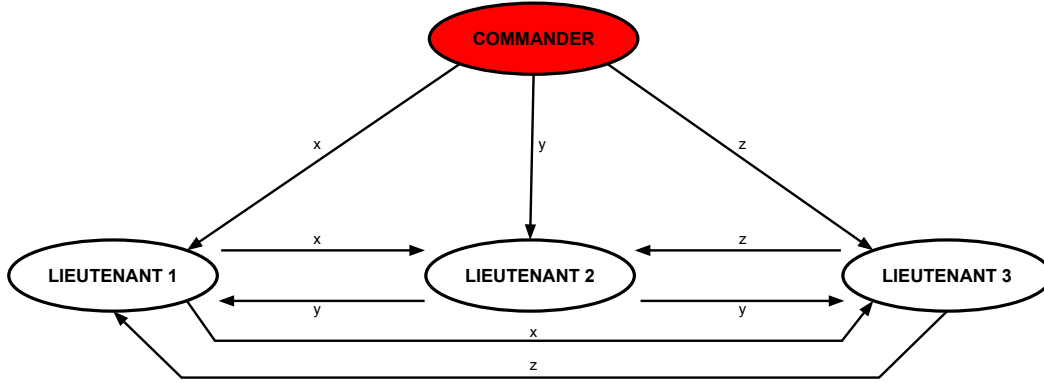


FIGURE 4. OM(1) with a disloyal lieutenant.

Before we can prove that the algorithm OM(m) works for an arbitrary m , we first need to prove the following lemma.

Lemma 1. *For any m and k , Algorithm OM(m) satisfies IC2 if there are more than $2k + m$ generals and at most k traitors.*

Proof. We will prove this by induction on m .

IC2 only cares when the commander is loyal. Using A1, we know that every message is delivered correctly. OM(0) works if the commander is loyal, since the lieutenants all receive the order. So it holds true for $m = 0$.

Now we assume it is true for $m - 1$, $m > 0$, and will prove for m .

In step 1, the loyal commander sends a value v to all $n - 1$ lieutenants. Each loyal lieutenant calls OM($m - 1$) with $n - 1$ generals. We know $n > 2k + 1$, so we know $n - 1 > 2k + m - 1$ or $n - 1 > 2k + (m - 1)$.

We use the inductive hypothesis and conclude that every loyal lieutenant gets $v_j = v$ for each loyal Lieutenant j . Since $n - 1 > 2k + (m - 1) \geq 2k$ and we know there is at most k traitors, the majority of $n - 1$ lieutenants are loyal. So every loyal lieutenant gets $v_i = v$ for a majority of the $n - 1$ values i . This gives him $\text{majority}(v_1, \dots, v_{n-1})$ in step 2. This proves IC2. \square

Theorem 2. *For any m , Algorithm OM(m) satisfies conditions IC1 and IC2 if there are more than $3m$ generals and at most m traitors.*

Proof. We prove this using induction on m .

It is simple to see that OM(0) satisfies IC1 and IC2 if there are no traitors. Assume the theorem is true for OM($m - 1$) and prove it for OM(m), when $m > 0$.

Case 1. The commander is loyal.

Use Lemma 1 and set k equal to m .

$$n > 2k + m$$

$$n > 3k.$$

We see that OM(m) satisfies IC2, and IC1 follows if the commander is loyal.

Case 2. The commander is a traitor.

We know there is at most m traitors, since one is the commander, there is at most $m - 1$ traitorous lieutenants. We also know that there are more than $3m$ generals, therefore there are more than $3m - 1$ lieutenants. $3m - 1 > 3(m - 1)$. We can apply the inductive hypothesis and see that OM($m - 1$) meets both IC1 and IC2. Since IC2 is satisfied, we know every loyal lieutenant receives the same set of values v_1, \dots, v_{n-1} . Therefore they come to the same value using the $\text{majority}(v_1, \dots, v_{n-1})$ function, proving IC1.

This proves that OM(m) satisfies IC1 and IC2 if there is more than $3m$ generals. \square

4. SIGNED MESSAGES

So far we have been looking at oral, unsigned messages. This makes finding a solution to the Byzantine Generals Problem more difficult because it allows the traitorous generals to lie. Restricting a traitor's ability to lie would make solving the problem easier. In this section we will explore that by having each general send signed messages with a signature that cannot be forged. To do this we need more assumptions:

- A4. (a) A loyal general's signature cannot be forged, and any alterations of the contents of their signed messages can be detected.
- (b) Anyone can verify the authenticity of a general's signature.

No assumptions are made about traitors' signature. It is possible for traitors to forge another traitor's signature and to collude in that manner.

This new algorithm is able to cope with m traitors for any amount of generals, due to the messages now being signed. The authors do note that it does not make much sense for less than $m + 2$ generals.

In this new algorithm, every when the commander sends an order he signs it. So each lieutenant receives an order with the commanders signature. They then add their signature to the message before sending it to the other lieutenants. Since each lieutenant will receive multiple messages with orders, we need a choice function to decide which order to follow. The only requirements for such a function are:

1. If the set V consists of the single element v , then $choice(V) = v$.
2. $choice(\emptyset) = \text{RETREAT}$, where \emptyset is the empty set.

The actual decision for deciding between different orders, as it will differ between implementations. It only matters that all loyal generals use the same decision making process.

In the algorithm we denote the value x signed by General i as $x : i$. That means $v : j : i$ would be the value v signed by General j and then signed by General i . Each General i maintains a set of properly signed orders V_i . A loyal commander should never have more than one element in V_i . V_i is specifically the set of orders a general has received, not the set of messages he has received.

Algorithm SM(m), m > 0

Initially $V_i = \{\}$

1. The commander signs and sends their value to every lieutenant.
2. For each i ,
 - a. If Lieutenant i receives a message from the commander of form $v : 0$ and they have not received any other order, then:
 - i. they set $V_i = \{v\}$.
 - ii. they send the message $v : 0 : i$ to every other lieutenant.
 - b. If Lieutenant i receives a message of the form $v : 0 : j_1 : \dots : j_k$ and v is not in V_i , then:
 - i. he adds v to V_i .
 - ii. if $k < m$, then he sends the message $v : 0 : j_1 : \dots : j_k : i$ to every other lieutenant, except for j_1, \dots, j_k .
 - c. For each i , when Lieutenant i receives no more messages, he follows the result from $choice(V_i)$.

5. MISSING COMMUNICATION PATHS

6. CONCLUSION

REFERENCES

1. L. Lamport, R. Shostak, and M. Pease, "The Byzantine Generals Problem", ACM Transactions on Programming Languages and Systems, Vol. 4, No. 3, 1982, 382-401.